

CDA5106 - Advanced Computer Architecture

Final Exam Review

Kobee Raveendran

1 Module 1: High-Performance Microprocessor Architecture

1.1 Module 1.2: Power Wall and Dennard Scaling

1.1.1 Notes

- energy: ability of a physical system to do work on other physical systems (unit: joule)
- power: rate at which energy is transformed (unit: watt; 1 watt = 1 joule delivered per second)
 - power = $V \cdot I$ (V = voltage, I = current)
- for capacitors:
 - energy stored = $0.5 \cdot C \cdot V^2$ (C = capacitance, V = voltage)
 - if a capacitor is drained at a frequency of f per second: power = $\frac{\text{energy}}{\text{second}} = 2 \cdot 0.5CV^2 = CV^2$
- Power wall problem
 - $P_{dyn} = ACV^2f$
 - A : fraction of gates actively switching
 - C : total capacitance of all gates
 - V : supply voltage
 - f : frequency of switching
- Power wall fundamentals
 - max frequency vs. threshold voltage:
 - $f_{max} = c \cdot \frac{(V-V_{thd})^{1.3}}{V}$
- Dennard Scaling Example (old)
 - if gate length (transistor size) scales by $S = 0.7$ (both length and width), then:
 - capacitance scales by $S = 0.7$
 - original area scales by $S^2 = 0.5$
 - number of transistors scales by $\frac{1}{S^2} \approx 2$
 - supply voltage (V) scales by $S = 0.7$
 - frequency (f) scales by $\frac{1}{S} = 1.4$
 - then, **dynamic power** $P_{dyn} = ACV^2f$
 - and **new dynamic power** $P'_{dyn} = A'C'V'^2f'$
 - $P'_{dyn} = (2A)(0.7C)(0.7V)^2(1.4f) \approx 1 \cdot ACV^2f = P_{dyn}$
- Post Dennard Scaling example (new)
 - capacitance scales by $S = 0.7$
 - number of transistors scales by $\frac{1}{S^2} = 2$

- supply voltage (V) cannot scale without also scaling threshold voltage (V_{thd}), and doing that increases static power exponentially
- frequency (f) scales by $\frac{1}{S} = 1.4$
- result: dynamic power doubles every generation
- $P_{dyn} = ACV^2f$
- $P'_{dyn} = A'C'V'^2f' = (2A)(0.7C)(1 \cdot V)^2(1.4f) \approx 2 \cdot P_{dyn}$

1.1.2 Exercises

- Suppose that instead of progressing at a ratio of 0.7, Moore's law slows down and transistor gate length scales at a ratio of 0.8 instead. Find the dynamic power consumption under *unlimited* and *limited* scaling for the next process generation.
 - Unlimited/old scaling rule
 - * gate length scales by $S = 0.8$
 - * capacitance scales by $S = 0.8$
 - * original area scales by $S^2 = 0.64$
 - * num transistors thus scales by $\frac{1}{S^2} = 1.56$
 - * supply voltage scales by $S = 0.8$
 - * frequency scales by $\frac{1}{S} = 1.25$
 - * dynamic power stays constant:

$$P_{dyn} = (1.56A)(0.8C)(0.8V)^2(1.25f)$$
 - leakage-limited/new scaling
 - * capacitance scales by $S = 0.8$
 - * num transistors scales by $\frac{1}{S^2} = 1.56$
 - * supply voltage does not scale without scaling threshold voltage too, which increases static power exponentially
 - * frequency scales by $\frac{1}{S}1.25$
 - * dynamic power consumption increases:

$$P'_{dyn} = (1.56A)(0.8C)(V)^2(1.25f) = 1.56 \cdot P_{dyn}$$
- With limited voltage scaling, suppose that we want to keep the dynamic power consumption constant in the next generation by keeping frequency constant and reduce die area. How much should we reduce die area to achieve that?
 - gate length scales by $S = 0.7$
 - capacitance scales by $S = 0.7$
 - original area scales by $S^2 = 0.5$
 - supply voltage and frequency are constant
 - dynamic power consumption must stay constant: $P'_{dyn} = P_{dyn}$

$$ACV^2f = A'(0.7C)V^2f \longrightarrow A = 0.7A'$$
 - number of transistors in the next generation: $A' = 1.4A$ (instead of $2A$ like before; i.e. 70% of $2A$)
 - thus die area shrinks by 30%
- Describe the difference between energy and power.
Power is the rate of energy consumption.
- Describe the impact of threshold voltage choice on static and dynamic power consumption as transistors are scaled down.
If threshold voltage is lowered, dynamic power decreases (nearly linearly) but static power increases exponentially.
- How has processor design adapted to the power wall problem?
Stalling frequency growth, multicore, and sophisticated power management (clock gating, voltage and frequency scaling, power gating).

1.1.3 Overview of ILP Techniques

Caches example

- processor with 1-ns clock
- 64KB cache memory with 2-ns read time, 95% hitrate
- 512MB main memory with 150-ns read time
- What is the average access time (AAT) in this memory system?

Answer:

- hits: $95 \cdot 2$ ns, misses: $5 \cdot (2 + 150)$ ns
- total = hit time + miss time = $190 + (10 + 750) = 950ns$
- $AAT = \frac{total}{100} = 9.5ns$

2 Module 2: Performance, Cost, and Reliability of Micro-processors

2.1 Performance Evaluation 1

2.1.1 Amdahl's Law

- performance improvement (“speedup”) is limited by the part you can’t improve
- (s) $Speedup_{enhanced}$ = best case speedup from gizmo alone
- (f) $Fraction_{enhanced}$ = fraction of task that gizmo can enhance
- $s_{overall} = \frac{1}{(1-f) + \frac{f}{s}}$

Example:

- jet plane wing simulation, where 1 run takes 1 week on your computer
- your program is 80% parallelizable
- new supercomputer has 100,000 processors
- $s = 100,000$
- $f = 0.8$
- overall speedup: $s_{overall} = \frac{1}{(1-f) + \frac{f}{s}} = \frac{1}{(1-0.8) + \frac{0.8}{100000}} \approx \frac{1}{0.2} = 5$
- only about 5 times faster (33 hours instead of 1 week), but not worth the high price tag (using a cheaper computer with only 100 processors instead yields a 4.8X speedup!)

More examples:

Ex 1:

- $f = 0.95$
- $s = 1.10$
- $s_{overall} = \frac{1}{(1-0.95) + \frac{0.95}{1.10}} = 1.094 \approx 1.10$

Ex 2:

- $f = 0.05$
- $s \rightarrow \infty$
- $s_{overall} = 1.053$

2.1.2 Run Time

- CPU time = clock cycle count \times cycle time
- cycles per instruction (CPI) = $\frac{\text{clock cycle count}}{\text{instruction count}}$
- CPU time = IC \times CPI \times CT

2.2 Performance Evaluation 2

Determine speedup by comparing program times with respect to a reference machine.

- arithmetic mean (which one should we trust?):

	Computer A	Computer B	B vs. A
Program P1	2X faster	4X faster	2X faster
Program P2	5X faster	15X faster	3X faster
Average	3.5X	9.5X	

Speedups:

- method 1: program-wise $\rightarrow \frac{2+3}{2} = 2.5\text{X}$ faster
- method 2: machine-wise $\rightarrow \frac{9.5}{3.5} = 2.71\text{X}$ faster

- geometric mean (consistent):

$$gmean = \sqrt[n]{\prod_{i=1}^n} = \exp\left(\frac{\frac{1}{n} \sum_{i=1}^n \ln(x_i)}{n}\right)$$

	Computer A	Computer B	B vs. A
Program P1	2X faster	4X faster	2X faster
Program P2	5X faster	15X faster	3X faster
Average	$\sqrt{10}$	$\sqrt{60}$	$\sqrt{6}$

Speedups:

- method 1: program-wise \rightarrow B is $\sqrt{2 \cdot 3} = \sqrt{6}\text{X}$ faster
- method 2: machine-wise \rightarrow B is $\sqrt{60} \cdot \sqrt{10} = \sqrt{6}\text{X}$ faster

- (also important): geometric standard deviation

$$gstdev = \exp\left(\sqrt{\frac{\prod_{i=1}^n (\ln x_i - \ln gmean)^2}{n}}\right)$$

in plain English: for each “component” speedup vs. ref machine, take its natural log and subtract the natural log of the gmean from that. Square it and multiply all of these together, then divide by n. Finally take the square root of this, then take e to the power of the result.

2.2.1 Exercises

Given the following table of speedups for machines A and B relative to a reference machine:

Prog	X (secs)	A (secs)	B (secs)
App 1	30	15	10
App 2	20	15	10
App 3	40	20	30
App 4	15	20	15

Compute the following (see post-computation table below to find them all):

- geometric speedup of machine A vs. base machine X
from the table, we find that A has a 1.41X speedup over X
- geometric speedup of machine B vs. base machine X
from the table, we find that B has a 1.68X speedup over X
- geometric speedup of machine B vs. machine A
from the table, we find that B has a 1.19X speedup over A

- geometric standard deviation of the speedup of machine A over machine X

$$gstd = \exp(\sqrt{\frac{1}{4} \cdot \ln^2(\frac{2}{1.41}) \ln^2(\frac{1.33}{1.41}) \ln^2(\frac{2}{1.41}) \ln^2(\frac{0.75}{1.41})})$$

$$gstd = 1.002255... \approx 1$$

Prog	A vs. X	B vs. X	B vs. A
App 1	2X	3X	1.5X
App 2	1.33X	2X	1.5X
App 3	2X	1.33X	0.67X
App 4	0.75X	1X	1.33X
Product	4X	8X	2X
gmean	1.41X	1.68X	1.19X

2.3 Cost and Reliability

2.3.1 Failure Rates (λ)

- λ = the number of failures that occur per unit time in a component/system
- FIT (failure in time) = number of failures in 10^9 hours
- example: 10,000 microprocessor chips used for 1,000 hours, and 8 of them fail. Failure rate is thus $\frac{8}{10,000 \cdot 1,000} = 8 \cdot 10^{-7}$ (failures per hour per chip) $\cdot 10^9$ hours = 800 FITs

2.3.2 Reliability Metrics

- $R(t)$ = probability that the system still works correctly at time t
- $W_N(t)$ = number of items (of the same kind) that would still be working at time t
- if λ is constant, then $R(t) = e^{-\lambda t}$
- Mean Time Between Failure (MTBF) = $\frac{1}{\lambda}$

2.3.3 System Reliability

Assume that:

- M components are in the system with failure rates $\lambda_1, \lambda_2, \dots, \lambda_m$
- for the system to work properly, all components must also work properly
- a component's reliability is independent of any other component's reliability
- then, system failure rate = sum of component's failure rates
- $R_{sys}(t) = R_1(t) \cdot R_2(t) \cdot \dots \cdot R_m(t) = e^{-\lambda_1 t} \dots e^{-\lambda_m t} = e^{-(\lambda_1 + \lambda_2 + \dots + \lambda_m)t} = e^{-\lambda_{sys} t}$

Other metrics:

- Mean Time To Repair (MTTR): mean time to repair/recover from a fault
- Mean Time Between Failure (MTBF): mean time between 2 consecutive failures
- if each failure is repaired, then $MTBF = MTTF + MTTR$
- usually, $MTTF \gg MTTR$, so MTBF and MTTF are often interchangeable

2.3.4 Examples

Assume a disk subsystem with:

- 10 disks each rated at 10^6 -hour MTTF
- 1 SCSI controller rated at $5 \cdot 10^5$ -hour MTTF
- 1 power supply rated at $2 \cdot 10^5$ -hour MTTF
- 1 fan rated at $2 \cdot 10^5$ -hour MTTF
- 1 SCSI cable rated at 10^6 -hour MTTF

Find the failure rate of the entire disk subsystem.
 $R_{sys}(t) = 10 \cdot \frac{1}{10^6} + \frac{1}{5 \cdot 10^5} + \frac{2}{2 \cdot 10^5} + \frac{1}{10^6} = \frac{10+2+5+5+1}{10^6} = \frac{23}{10^6} = \frac{23,000}{10^9} = 23,000 \text{ FIT}$
Thus, MTTF = $\frac{1}{\lambda_{sys}} = \frac{1}{23,000} \cdot 10^9 \approx 43,500 \text{ hours}$

3 Instruction Set Design

3.1 Instruction Set Architecture 1

3.1.1 Styles of ISAs

- Stack:
 - push <addr>, pop <addr> (or ALU instructions)
 - ALU: pop two entries, perform ALU operation, push result onto stack
 - compact instruction format
 - * all calculation operations take 0 operands
 - * flexible; used for compiling Java bytecode (not dependent on registers in architecture [but all have stacks])
- Accumulator:
 - load/store/ALU <addr> (result affects accumulator register)
 - also very compact (all operations take 1 operand, the other is implicitly the accumulator register)
 - less dependence on memory than stack-based
- Register-memory:
 - load/store/ALU <reg>, <reg/addr>
 - * at most one operand can be a memory address
 - * leftmost register is the destination (if applicable)
- Load-Store:
 - load/store <reg>, <addr>
 - ex: ALU <reg1>, <reg2>, <reg3>: reg1 is destination, other 2 are source registers

Example for adding two numbers:

Stack	Accumulator	Register-memory	Load-Store
push A	load A	load R1 A	load R1 A
push B	add B	add R1 B	load R2 B
add	store C	store C R1	add R3 R1 R2
pop C			store C R3

Pros/cons:

- load-store
 - (+) fixed length instructions possible (allows for easy fetch/decode)
 - (+) simpler hardware: efficient pipeline and potentially lower cycle time
 - (-) higher instruction count

- (-) fixed-length instructions can be wasteful (more bits than needed for some instructions)
- register-memory
 - (+) no need for extra loads
 - (+) better usage of bits (these pros lead to better code density)
 - (-) destroys source operand(s) (i.e. `add R1 R2`)
 - (-) may impact cycles per instruction
- memory-memory
 - (+) most compact (code density)
 - (-) high memory traffic (thus bottlenecked by memory)

3.2 Instruction Set Architecture 2

3.2.1 RISC and Common Addressing Modes

- register
 - `add R4 R3 // R4 = R4 + R3`
 - used when value is in a register
- immediate
 - `add R4 #3 // R4 = R4 + 3`
 - used for small constants (which occur frequently)
- displacement
 - `add R4 100(R1) // R4 = R4 + MEM[100 + R1]`
 - accesses the frame (arguments, local variables)
 - access the global data segment
 - accesses the fields of a data struct
- register deferred/register indirect
 - `add R3 (R1) // R3 = R3 + MEM[R1]`
 - accesses using a computed memory address
- indexed
 - `add R3 (R1 + R2) // R3 = R3 + MEM[R1 + R2]`
 - array accesses: $R1 = \text{base}$, $R2 = \text{index}$
- direct/absolute
 - `add R1 (1001) // R1 = R1 + MEM[1001]`
 - accessing global (“static”) data
- memory deferred/memory indirect
 - `add R1 @(R3) // R1 = R1 + MEM[MEM[R3]]`
 - pointer dereferencing: `x = *p`; (if `p` is not register-allocated)
- autoincrement/postdecrement
 - `add R1 (R2)+ // R1 = R1 + MEM[R2]; R2 = R2 + d` (d is the size of the operation)
 - looping through arrays, stack pop
- autodecrement/predecrement
 - `add R1 -(R2) // R1 = R1 + MEM[R2]; R2 = R2 - d` (d is the size of the operation)

- same uses as autoincrement, stack push
- scaled
 - `add R1 100(R2) [R3] // R1 = R1 + MEM[100 + R2 + R3 * d]`
 - array accesses for non-byte-sized elements

3.3 Instruction Set Architecture 3

3.3.1 Condition Codes (for branch instructions)

- Z
 - zero flag
 - indicates the result of an arithmetic/logical expression is zero
- C
 - carry flag
 - indicates that an operation has a carry out. Enables numbers larger than a single word to be added/subtracted
- S/N
 - sign/negative flag
 - indicates the result of an operation is negative
- V/O/W
 - overflow flag
 - indicates the result of an operation is too large to fit in a register (using 2's complement representation)

3.3.2 Instruction Encoding Tradeoffs

- variable width
 - (+) very versatile, uses memory efficiently
 - (-) instruction words must be decoded before number of bytes is known (harder to fetch/decode)
- fixed width
 - (+) every instruction word is an instruction, thus easier to fetch/decode
 - (-) uses memory inefficiently (same num. bits even for short instructions)
- hybrid
 - primarily for embedded processors to conserve memory
 - often use a subset of instructions, fewer registers, or even instruction compression

3.4 ISA Examples

3.4.1 MIPS

Characteristics:

- load/store ISA: only loads/stores can have memory operands, makes for easy pipelining and uniform instruction width
- fixed instruction width = easy fetch and decode
- small number of addressing modes = easy pipelining
- large register file: 32 integer and 32 floating point registers
- aligned memory = easy data fetching

- quantitatively designed

Instruction format:

- R: `op`, `rs`, `rt`, `rd`, `shamt`, `funct` (`shamt` = shift amount, `funct` = ALU function)
- I: `op`, `rs`, `rt`, 16-bit address
- J: `op`, 26-bit address